

# Process Synchronization

STDISCM S17

Atasha Dominique C. Pidlaon

# Problem Overview

## WHAT WE WANT TO SOLVE

**Objective:** Manage dungeon queuing efficiently with process synchronization

### Constraints:

- Max n concurrent dungeon instances
- Parties must contain 1 Tank, 1 Healer, 3 DPS
- Avoid deadlock & starvation
- Random dungeon duration t1 - t2 seconds

**Solution:** Implement thread synchronization for fairness & concurrency

**Deadlock** is a state where two or more processes are **permanently waiting** for each other to release a locked resource, causing the entire system to halt.

**Potential Scenario:**

- A dungeon instance successfully locks initial player roles (e.g., **Tank** and **DPS** players).
- It then waits indefinitely for a missing role (**Healer**), holding the locked Tank/DPS resources hostage.
- These locked resources are unavailable to any other waiting party, leading to **system starvation**.

**Solution:** Ensure resource availability before forming a party

# Deadlock Explanation

# Starvation

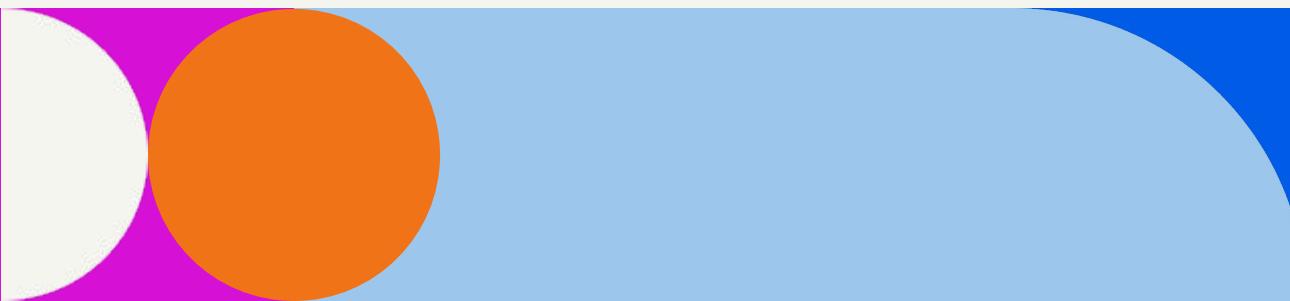
**Starvation** is when a ready player (process) is **repeatedly ignored** and cannot access needed resources (party slot), even if the resource is available.

## Potential Scenario

- The system constantly forms new parties, but if the logic **always favors fresh arrivals** or certain role combinations, older players are **perpetually overlooked** in the queue.
- The ignored player is ready but **never selected** for a party, forcing them to wait indefinitely.

**Strategy:** Implement a **Fair Queuing Mechanism**.

**Result:** The assignment logic ensures **all waiting players are considered equally** and will eventually be assigned a party, eliminating the risk of indefinite waiting.



# Synchronization Mechanisms Used

## MUTEX

- Prevents **race conditions** on shared data.
- Ensures **only one thread** modifies critical data (like player counts) at a time.
- Guarantees **data integrity**.

## THREADING

- Enables **concurrent handling** of all dungeon instances.
- Each dungeon runs on its own **separate thread**.
- Maximizes system **throughput**.

# Code Implementation

## Mutex Usage:

```
// --- Concurrency and Synchronization ---
mutex coutMutex;           // For synchronized console output
mutex statsMutex;          // For dungeon statistics
mutex playerMutex;         // For protecting player counters
condition_variable cv;     // For signaling player availability
```

```
while (true) {
    unique_lock<mutex> lock(playerMutex);
```

## Threading for Dungeon Instances:

```
instanceStats.resize(numDungeonsToRun);
vector<thread> dungeonThreads;

cout << "\nStarting dungeon instances..." << endl;

for (int i = 0; i < numDungeonsToRun; i++) {
    dungeonThreads.emplace_back(dungeonConsumer, i);
}
```

## Randomized Dungeon Time Simulation:

```
// --- Simulate Dungeon Run ---
{
    lock_guard<mutex> statsLock(statsMutex);
    instanceStats[instanceId].active = true;
}
{
    lock_guard<mutex> coutLock(coutMutex);
    cout << "\nQueueing up players for Dungeon Instance " << instanceId + 1 << endl;
    printDungeonStatuses();
}

int dungeonTime = dis(gen);
this_thread::sleep_for(chrono::seconds(dungeonTime));
```